

CSE 215:SQL

Dr. Mohammed Eunus Ali
(eunus@cse.buet.ac.bd)

Department of Computer Science and Engineering
Bangladesh University of Engineering and Technology (BUET)
Dhaka-1000, Bangladesh

(with some slides integrated from those of Jenifer Widom, Alon Halevy, Carlo Curino, and Michael Stonebraker.)

SQL Introduction

- Standard language for querying and manipulating data
Structured Query Language (SQL)
- Many standards out there: **SQL92, SQL2, SQL3, SQL99**
- **Vendors support various subsets of these**

SQL

Basic form: (many many more bells and whistles in addition)

Select attributes

From relations (possibly multiple, joined)

Where conditions (selections)

Selections

Company(sticker, name, country, stockPrice)

Find all US companies whose stock is > 50:

```
SELECT *  
FROM Company  
WHERE country="USA" AND stockPrice > 50
```

Output schema: R(sticker, name, country, stockPrice)

Selections

- **What you can use in WHERE:**

attribute names of the relation(s) used in the FROM.

comparison operators: =, <>, <, >, <=, >=

apply arithmetic operations: stockprice*2

operations on strings (e.g., "||" for concatenation).

Lexicographic order on strings.

Pattern matching: s LIKE p

Special stuff for comparing dates and times.

Schema

- **Product (pname, price, category, maker)**
- **Purchase (buyer, seller, store, product)**
- **Company (cname, stockPrice, country)**
- **Person(pname, phoneNumber, city)**

Projections

- Select only a subset of the attributes

```
SELECT name, stockPrice  
FROM Company  
WHERE country="USA" AND stockPrice > 50
```

- Input schema: **Company(sticker, name, country, stockPrice)**
- Output schema: **R(name, stock price)**

Projections

Rename the attributes in the resulting table

```
SELECT name AS company, stockprice AS price  
FROM Company  
WHERE country="USA" AND stockPrice > 50
```

Input schema: **Company**(sticker, name, country, stockPrice)

Output schema: **R**(company, price)

Joins

Find names of people living in Seattle that bought gizmo products, and the names of the stores they bought from

```
SELECT pname, store
FROM Person, Purchase
WHERE pname=buyer AND city="Seattle"
      AND product="gizmo"
```

Disambiguating Attributes

Find names of people buying telephony products:

Product (name, price, category, maker)

Purchase (buyer, seller, store, product)

Person(name, phoneNumber, city)

```
SELECT Person.name
FROM Person, Purchase, Product
WHERE Person.name=Purchase.buyer
      AND Person.name=Product.name
      AND Product.category="telephony"
```

Tuple Variables

Find pairs of companies making products in the same category

```
SELECT product1.maker, product2.maker
FROM   Product AS product1, Product AS product2
WHERE  product1.category=product2.category
      AND product1.maker <> product2.maker
```

Product (name, price, category, maker)

Tuple Variables

Tuple variables introduced automatically by the system:

Product (name, price, category, maker)

```
SELECT name  
FROM   Product  
WHERE  price > 100
```

Becomes:

```
SELECT Product.name  
FROM   Product AS Product  
WHERE  Product.price > 100
```

Doesn't work when Product occurs more than once:
In that case the user needs to define variables explicitly.

Meaning (Semantics) of SQL Queries

SELECT a1, a2, ..., ak
FROM R1 AS x1, R2 AS x2, ..., Rn AS xn
WHERE Conditions

1. Nested loops:

```
Answer = {}  
for x1 in R1 do  
    for x2 in R2 do  
        .....  
        for xn in Rn do  
            if Conditions  
                then Answer = Answer U {(a1,...,ak)}  
return Answer
```

Meaning (Semantics) of SQL Queries

SELECT a1, a2, ..., ak
FROM R1 AS x1, R2 AS x2, ..., Rn AS xn
WHERE Conditions

2. Parallel assignment

```
Answer = {}  
for all assignments x1 in R1, ..., xn in Rn do  
    if Conditions then Answer = Answer U {(a1,...,ak)}  
return Answer
```

Doesn't impose any order !
Like Datalog

Meaning (Semantics) of SQL Queries

SELECT a1, a2, ..., ak
FROM R1 AS x1, R2 AS x2, ..., Rn AS xn
WHERE Conditions

3. Translation to Datalog: one rule

$\text{Answer}(a_1, \dots, a_k) \leftarrow R_1(x_{11}, \dots, x_{1p}), \dots, R_n(x_{n1}, \dots, x_{np}), \text{Conditions}$
--

Meaning (Semantics) of SQL Queries

SELECT a1, a2, ..., ak
FROM R1 AS x1, R2 AS x2, ..., Rn AS xn
WHERE Conditions

4. Translation to Relational algebra:

$\Pi_{a1, \dots, ak} (\sigma_{\text{Conditions}} (R1 \times R2 \times \dots \times Rn))$

Select-From-Where queries are precisely Select-Project-Join

Movie Database Schema

Schema

Movie (title, year, length, inColor, studioName, producerC#)
StarsIn (movieTitle ,movieYear, starName)
MovieStar(name, address, gender, birthdate)
MovieExec(name, address, cert#, netWorth)
Studio(name, address, presC#)

Ordering the Output

Show all Disney movies of 1990 in ascending order of their lengths.

Movie (title, year, length, inColor, studioName, producerC#)

```
SELECT name, address, length
FROM    Movie
WHERE   studioName = 'Disney' AND year=1990
ORDER BY length;
```

If lengths are equal then in alphabetical order.

Union, Intersection, Difference

Find the name and address of all female movie stars who are
Also movie executives with a net worth over \$100,000,000.

```
(SELECT name, address  
FROM MovieStar  
WHERE gender="F")
```

INTERSECTION

```
(SELECT name, address  
FROM MovieExec  
WHERE networth > 100,000,000
```

You must have the same attribute names (otherwise: rename).

Union, Intersection, Difference

Find the names and addresses of movie stars who are movie executives.

```
(SELECT name, address  
FROM MovieStar)
```

EXCEPT

```
(SELECT name, address  
FROM MovieExec)
```

Union, Intersection, Difference

Find the titles and years of movies that appered in either the Movie or StarsIn

```
(SELECT title, year  
FROM Movie)
```

EXCEPT

```
(SELECT movieTtile As title, movieYear as year  
FROM StarsIn)
```

You must have the same attribute names (otherwise: rename).

Subqueries

- A query that is a part of another query is called subquery.
- Subqueries can be used in WHERE or FROM clauses

E.g., Show the producer name of Star Wars:

```
SELECT name
FROM   MovieExec
WHERE  cert# =
      (SELECT producerC#
       FROM   Movie
       WHERE  title = 'Star Wars');
```

In this case, the subquery returns one value.
If it returns more, it's a **run-time error**.

Can say the same thing without a subquery:

```
SELECT name  
FROM    Movie, MovieExec  
WHERE title = 'Star Wars' AND producerC# = cert#
```

Is this query equivalent to the previous one ?

Subqueries Returning Relations

Find studioname of Harrison Ford Movie.

```
SELECT studioname
FROM Movie
WHERE (name,year) IN
      (SELECT movieTitle, movieYear
       FROM StartsIn
       WHERE starName = "Harrison Ford");
```

Here the subquery returns a set of values.

Find the producer names of Harrison Fords' movies.

Subqueries Returning Relations

You can also use:

$s \text{ IN } R$

$s \text{ NOT IN } R$

$s > \text{ALL } R$

$s > \text{ANY } R$

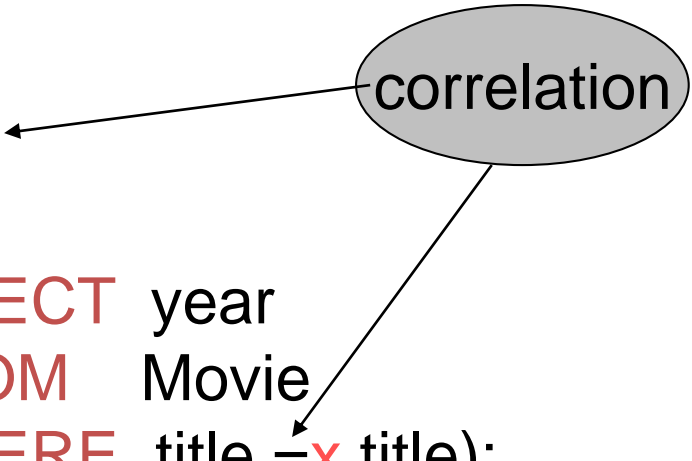
$\text{EXISTS } R$

Correlated Queries

Movie (title, year, director, length)

Find movies whose title appears more than once.

```
SELECT title
FROM Movie AS x
WHERE year < ANY
      (SELECT year
       FROM Movie
       WHERE title = x.title);
```



The diagram illustrates the correlation in the query. A gray oval labeled "correlation" has two arrows. One arrow points to the alias 'x' in the 'FROM Movie AS x' clause. The other arrow points to the 'x' in the 'WHERE title = x.title' clause of the subquery, indicating that the subquery is correlated with the outer query's alias 'x'.

Note (1) scope of variables (2) this can still be expressed as single SFW

Subqueries in FORM clause

```
SELECT studioName
FROM Movie
WHERE (name,year) IN
      (SELECT movieTitle, movieYear
       FROM StartsIn
       WHERE starName = "Harrison Ford");
```

Removing Duplicates

Find the producer names of Harrison Fords' movies (one producer name should appear only once.)

```
SELECT name
FROM   MovieExec
WHERE  cert# IN
      (SELECT producerC#
       FROM   Movie
       WHERE  (name,year) IN
            (SELECT movieTitle, movieYear
             FROM   StartsIn
             WHERE  starName = "Harrison Ford");
```

Aggregation

Schema:

Product (pname, price, category, maker)

Purchase (buyer, seller, store, product)

Company (cname, stockPrice, country)

Person(pname, phoneNumber, city)

```
SELECT Sum(price)
FROM Product
WHERE maker="Toyota"
```

SQL supports several aggregation operations:

SUM, MIN, MAX, AVG, COUNT

Aggregation: Count

```
SELECT Count(*)  
FROM Product  
WHERE year > 1995
```

Except COUNT, all aggregations apply to a single attribute

Aggregation: Count

COUNT applies to duplicates, unless otherwise stated:

```
SELECT Count(name, category)      same as Count(*)  
FROM Product  
WHERE year > 1995
```

Better:

```
SELECT Count(DISTINCT name, category)  
FROM Product  
WHERE year > 1995
```

Simple Aggregation

Purchase(product, date, price, quantity)

Example 1: **find total sales for the entire database**

```
SELECT Sum(price * quantity)
FROM Purchase
```

Example 1': **find total sales of bagels**

```
SELECT Sum(price * quantity)
FROM Purchase
WHERE product = 'bagel'
```


Simple Aggregations

Product	Date	Price	Quantity
Bagel	10/21	0.85	15
Banana	10/22	0.52	7
Banana	10/19	0.52	17
Bagel	10/20	0.85	20

Grouping and Aggregation

Usually, we want aggregations on certain parts of the relation.

Purchase(product, date, price, quantity)

Example 2: **find total sales after 9/1 per product.**

```
SELECT    product, Sum(price*quantity) AS TotalSales
FROM      Purchase
WHERE     date > "9/1"
GROUPBY   product
```

Grouping and Aggregation

1. Compute the relation (I.e., the FROM and WHERE).
2. Group by the attributes in the GROUPBY
3. Select one tuple for every group (and apply aggregation)

SELECT can have (1) grouped attributes or (2) aggregates.

First compute the relation (date > “9/1”)
then group by product:

Product	Date	Price	Quantity
Banana	10/19	0.52	17
Banana	10/22	0.52	7
Bagel	10/20	0.85	20
Bagel	10/21	0.85	15

Then, aggregate

Product	TotalSales
Bagel	\$29.75
Banana	\$12.48

```
SELECT    product, Sum(price*quantity) AS TotalSales
FROM      Purchase
WHERE     date > "9/1"
GROUPBY   product
```

Another Example

Product	SumSales	MaxQuantity
Banana	\$12.48	17
Bagel	\$29.75	20

For every product, what is the total sales and max quantity sold?

```
SELECT      product, Sum(price * quantity) AS  
SumSales  
                                     Max(quantity) AS MaxQuantity  
FROM        Purchase  
GROUP BY   product
```

HAVING Clause

Same query, except that we consider only products that had at least 100 buyers.

```
SELECT    product, Sum(price * quantity)
FROM      Purchase
WHERE     date > "9/1"
GROUP BY  product
HAVING    Sum(quantity) > 30
```

HAVING clause contains conditions on aggregates.

General form of Grouping and Aggregation

SELECT **S**
FROM R_1, \dots, R_n
WHERE **C1**
GROUP BY a_1, \dots, a_k
HAVING **C2**

S = may contain attributes a_1, \dots, a_k and/or any aggregates but NO OTHER ATTRIBUTES

C1 = is any condition on the attributes in R_1, \dots, R_n

C2 = is any condition on aggregate expressions

General form of Grouping and Aggregation

```
SELECT    S  
FROM      R1,...,Rn  
WHERE     C1  
GROUP BY a1,...,ak  
HAVING    C2
```

Evaluation steps:

1. **Compute the FROM-WHERE part, obtain a table with all attributes in R_1, \dots, R_n**
2. **Group by the attributes a_1, \dots, a_k**
3. **Compute the aggregates in C2 and keep only groups satisfying C2**
4. **Compute aggregates in S and return the result**

Aggregation

Author(login,name)

Document(url, title)

Wrote(login,url)

Mentions(url,word)

- Find all authors who wrote at least 10 documents:

```
Select author.name  
From author, wrote  
Where author.login=wrote.login  
Groupby author.name  
Having count(wrote.url) > 10
```

- **Find all authors who have a vocabulary over 10000:**

```
Select author.name  
From author, wrote, mentions  
Where author.login=wrote.login and wrote.url=mentions.url  
Groupby author.name  
Having count(distinct mentions.word) > 10000
```

Exercises

Product (pname, price, category, maker)

Purchase (buyer, seller, store, product)

Company (cname, stock price, country)

Person(per-name, phone number, city)

Ex #1: Find people who bought telephony products.

Ex #2: Find names of people who bought American products

Ex #3: Find names of people who bought American products and did not buy French products

Ex #4: How much money did Fred spend on purchases?

Ex #5: What is the number and sum of the product sales by country of origin?

Modifying the Database

Three kinds of modifications

- **Insertions**
- **Deletions**
- **Updates**

Sometimes they are all called “updates”

Insertions

General form:

```
INSERT INTO R(A1,..., An) VALUES (v1,..., vn)
```

Example: Insert a new purchase to the database:

```
INSERT INTO Purchase(buyer, seller, product, store)  
VALUES ('Joe', 'Fred', 'wakeup-clock-espresso-machine',  
        'The Sharper Image')
```

Missing attribute → NULL.

May drop attribute names if give them in order.

Insertions

```
INSERT INTO PRODUCT(name)  
  
    SELECT DISTINCT Purchase.product  
    FROM     Purchase  
    WHERE    Purchase.date > "10/26/01"
```

The query replaces the VALUES keyword.
Here we insert *many* tuples into PRODUCT

Insertion: an Example

Product(name, listPrice, category)
Purchase(prodName, buyerName, price)

prodName is foreign key in Product.name

Suppose database got corrupted and we need to fix it:

Product

name	listPrice	category
gizmo	100	gadgets

Purchase

prodName	buyerName	price
camera	John	200
gizmo	Smith	80
camera	Smith	225

Task: insert in Product all prodNames from Purchase

Insertion: an Example

```
INSERT INTO Product(name)
```

```
SELECT DISTINCT prodName
```

```
FROM Purchase
```

```
WHERE prodName NOT IN (SELECT name FROM Product
```

name	listPrice	category
gizmo	100	Gadgets
camera	-	-

Insertion: an Example

```
INSERT INTO Product(name, listPrice)
```

```
SELECT DISTINCT prodName, price
```

```
FROM Purchase
```

```
WHERE prodName NOT IN (SELECT name FROM Product)
```

name	listPrice	category
gizmo	100	Gadgets
camera	200	-
camera ??	225 ??	-

← Depends on the implementation

Deletions

Example:

```
DELETE FROM PURCHASE  
WHERE seller = 'Joe' AND  
product = 'Brooklyn Bridge'
```

Factoid about SQL: there is no way to delete only a single occurrence of a tuple that appears twice in a relation.

Updates

Example:

```
UPDATE PRODUCT
SET price = price/2
WHERE Product.name IN
      (SELECT product
       FROM Purchase
       WHERE Date ='Oct, 25, 1999');
```

Data Definition in SQL

So far we have seen the ***Data Manipulation Language***, DML
Next: ***Data Definition Language*** (DDL)

Data types:

Defines the types.

Data definition: defining the schema.

- Create tables
- Delete tables
- Modify table schema

Indexes: to improve performance

Data Types in SQL

- Character strings (fixed or varying length)
- Bit strings (fixed or varying length)
- Integer (SHORTINT)
- Floating point
- Dates and times

Domains (=types) will be used in table declarations.

To reuse domains:

```
CREATE DOMAIN address AS VARCHAR(55)
```

Creating Tables

Example:

```
CREATE TABLE Person(  
    name VARCHAR(30),  
    social-security-number INTEGER,  
    age SHORTINT,  
    city VARCHAR(30),  
    gender BIT(1),  
    Birthdate DATE  
);
```


Deleting or Modifying a Table

Deleting:

Example: `DROP Person;`

Altering: (adding or removing an attribute).

Example:

```
ALTER TABLE Person  
    ADD phone CHAR(16);  
  
ALTER TABLE Person  
    DROP age;
```

What happens when you make changes to the schema?

Default Values

Specifying default values:

```
CREATE TABLE Person(  
    name                VARCHAR(30),  
    social-security-number INTEGER,  
    age                 SHORTINT DEFAULT 100,  
    city                VARCHAR(30) DEFAULT 'Seattle',  
    gender              CHAR(1)  DEFAULT '?',  
    Birthdate           DATE
```

The default of defaults: NULL

Indexes

REALLY important to speed up query processing time.


Suppose we have a relation

Person (name, age, city)

```
SELECT *  
FROM   Person  
WHERE  name = "Smith"
```

Sequential scan of the file Person may take long

- **Create an index on name:**



Adam	Betty	Charles	Smith
------	-------	---------	------	-------	------

- **B+ trees have fan-out of 100s: max 4 levels**
!

Creating Indexes

Syntax:

```
CREATE INDEX nameIndex ON Person(name)
```

Creating Indexes

Indexes can be created on more than one attribute:

Example:

```
CREATE INDEX doubleindex ON  
Person (age, city)
```

Helps in:

```
SELECT *  
FROM Person  
WHERE age = 55 AND city = "Seattle"
```

But not in:

```
SELECT *  
FROM Person  
WHERE city = "Seattle"
```

Creating Indexes

Indexes can be useful in range queries too:

```
CREATE INDEX ageIndex ON Person (age)
```

B+ trees help in:

```
SELECT *  
FROM Person  
WHERE age > 25 AND age < 28
```

Why not create indexes on everything?

Defining Views

Views are relations, except that they are not physically stored.

For presenting different information to different users

Employee(ssn, name, department, project, salary)

```
CREATE VIEW Developers AS  
  SELECT name, project  
  FROM Employee  
  WHERE department = "Development"
```

Payroll has access to **Employee**, others only to **Developers**

A Different View

Person(name, city)

Purchase(buyer, seller, product, store)

Product(name, maker, category)

```
CREATE VIEW Seattle-view AS
```

```
SELECT buyer, seller, product, store
```

```
FROM Person, Purchase
```

```
WHERE Person.city = "Seattle" AND  
       Person.name = Purchase.buyer
```

We have a new virtual table:

Seattle-view(buyer, seller, product, store)

A Different View

We can later use the view:

```
SELECT name, store
FROM Seattle-view, Product
WHERE Seattle-view.product = Product.name AND
      Product.category = "shoes"
```

What Happens When We Query a View ?

```
SELECT name, Seattle-view.store
FROM   Seattle-view, Product
WHERE  Seattle-view.product = Product.name AND
       Product.category = "shoes"
```



```
SELECT name, Purchase.store
FROM   Person, Purchase, Product
WHERE  Person.city = "Seattle" AND
       Person.name = Purchase.buyer AND
       Purchase.product = Product.name AND
       Product.category = "shoes"
```

Types of Views

- **Virtual views:**
 - Used in databases
 - Computed only on-demand – slow at runtime
 - Always up to date
- **Materialized views**
 - Used in data warehouses (but recently also in DBMS)
 - Precomputed offline – fast at runtime
 - May have stale data

Updating Views

How can I insert a tuple into a table that doesn't exist?

Employee(ssn, name, department, project, salary)

```
CREATE VIEW Developers AS  
  SELECT name, project  
  FROM Employee  
  WHERE department = "Development"
```

If we make the
following insertion:

```
INSERT INTO Developers  
VALUES("Joe", "Optimizer")
```

It becomes:

```
INSERT INTO Employee  
VALUES(NULL, "Joe", NULL, "Optimizer", NULL)
```

Non-Updatable Views

```
CREATE VIEW Seattle-view AS  
  
    SELECT seller, product, store  
    FROM Person, Purchase  
    WHERE Person.city = "Seattle" AND  
           Person.name = Purchase.buyer
```

How can we add the following tuple to the view?

("Joe", "Shoe Model 12345", "Nine West")

What do we put in the Person.name and Purchase.buyer columns?

Answering Queries Using Views

- **What if we want to *use* a set of views to answer a query.**
- **Why?**
 - The obvious reason...
 - Answering queries over web data sources.
- ***Very cool stuff! (i.e., I did a lot of research on this).***

Reusing a Materialized View

- Suppose I have **only** the result of **SeattleView**:

```
SELECT buyer, seller, product, store
FROM    Person, Purchase
WHERE    Person.city = 'Seattle'  AND
           Person.per-name = Purchase.buyer
```

- and I want to answer the query

```
SELECT buyer, seller
FROM    Person, Purchase
WHERE    Person.city = 'Seattle'  AND
           Person.per-name = Purchase.buyer

AND

           Purchase.product='gizmo'.
```

Then, I can rewrite the query using the view.

Query Rewriting Using Views

Rewritten query:

```
SELECT buyer, seller  
FROM    SeattleView  
WHERE   product= 'gizmo'
```

Original query:

```
SELECT buyer, seller  
FROM    Person, Purchase  
WHERE   Person.city = 'Seattle'  AND  
          Person.per-name = Purchase.buyer  
AND  
          Purchase.product='gizmo'.
```

Another Example

- I still have **only** the result of SeattleView:

```
SELECT buyer, seller, product, store
FROM Person, Purchase
WHERE Person.city = 'Seattle' AND
      Person.per-name = Purchase.buyer
```

- but I want to answer the query

```
SELECT buyer, seller
FROM Person, Purchase
WHERE Person.city = 'Seattle' AND
      Person.per-name = Purchase.buyer
```

AND

```
Person.Phone LIKE '206 543 %'.
```

And Now?

- I still have **only** the result of **SeattleView**:

```
SELECT buyer, seller, product, store
FROM    Person, Purchase
WHERE    Person.city = 'Seattle'  AND
           Person.per-name = Purchase.buyer
```

- but I want to answer the query

```
SELECT buyer, seller
FROM    Person, Purchase, Product
WHERE    Person.city = 'Seattle'  AND
           Person.per-name = Purchase.buyer
```

AND

```
Person.Phone LIKE '206 543 %' AND
Purchase.product = Product.name.
```

And Now?

- I still have **only** the result of:
SELECT seller, buyer, Sum(Price)
FROM Purchase
WHERE Purchase.store = 'The Bon'
Group By seller, buyer
- but I want to answer the query
SELECT seller, Sum(Price)
FROM Purchase
WHERE Person.store = 'The Bon'
Group By seller

And what if it's the other way around?

Finally...

- I still have **only** the result of:

```
SELECT seller, buyer, Count(*)  
FROM Purchase  
WHERE Purchase.store = 'The Bon'  
Group By seller, buyer
```

- but I want to answer the query

```
SELECT seller, Count(*)  
FROM Purchase  
WHERE Person.store = 'The Bon'  
Group By seller
```

The General Problem

- **Given a set of views V_1, \dots, V_n , and a query Q , can we answer Q using only the answers to V_1, \dots, V_n ?**
- **Why do we care?**
 - We can answer queries more efficiently.
 - We can query data sources on the WWW in a principled manner.
- **Many, many papers on this problem.**
- **The best performing algorithm: The MiniCon Algorithm, (Pottinger & (Ha)Levy, 2000).**

Querying the WWW

- Assume a virtual schema of the WWW, e.g.,
 - Course(number, university, title, prof, quarter)
- Every data source on the web contains the answer to a view over the virtual schema:

UW database: **SELECT** number, title, prof
 FROM Course
 WHERE univ='UW' AND
 quarter='2/02'

Stanford database: **SELECT** number, title, prof,
 quarter
 FROM Course
 WHERE univ='Stanford'

User query: find all professors who teach “database systems”

Null Values and Outerjoins

- If $x = \text{Null}$ then $4 \cdot (3 - x) / 7$ is still NULL
- If $x = \text{Null}$ then $x = \text{"Joe"}$ is UNKNOWN
- Three boolean values:
 - FALSE = 0
 - UNKNOWN = 0.5
 - TRUE = 1

M.W.: NULL is not equal to NULL in SQL.

Null Values and Outerjoins

- **C1 AND C2 = min(C1, C2)**
- **C1 OR C2 = max(C1, C2)**
- **NOT C1 = 1 - C1**

```
SELECT *  
FROM Person  
WHERE (age < 25) AND  
        (height > 6 OR weight > 190)
```

**Rule in SQL: include only tuples that yield
TRUE**

Null Values and Outerjoins

Unexpected behavior:

```
SELECT *  
FROM   Person  
WHERE  age < 25 OR age >= 25
```

Some Persons are not included !

Null Values and Outerjoins

Can test for NULL explicitly:

- x IS NULL
- x IS NOT NULL

SELECT *

FROM Person

**WHERE age < 25 OR age >= 25 OR age IS
NULL**

Now it includes all Persons

Null Values and Outerjoins

Explicit joins in SQL:

Product(name, category)
Purchase(prodName, store)

```
SELECT Product.name, Purchase.store  
FROM    Product JOIN Purchase ON  
          Product.name = Purchase.prodName
```

Same as:

```
SELECT Product.name, Purchase.store  
FROM    Product, Purchase  
WHERE    Product.name = Purchase.prodName
```

But Products that never sold will be lost !

Null Values and Outerjoins

Left outer joins in SQL:

Product(name, category)

Purchase(prodName, store)

SELECT Product.name, Purchase.store

FROM Product **LEFT OUTER JOIN** Purchase **ON**
Product.name = Purchase.prodName

Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

Name	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz
OneClick	-

Outer Joins

- **Left outer join:**
 - Include the left tuple even if there's no match
- **Right outer join:**
 - Include the right tuple even if there's no match
- **Full outer join:**
 - Include the both left and right tuples even if there's no match

SQL: Constraints and Triggers

- Chapter 6 Ullman and Widom
- Certain properties we'd like our database to hold
- Modification of the database may break these properties
- Build handlers into the database definition
- **Key constraints**
- **Referential integrity constraints.**

Declaring a Primary Keys in SQL

```
CREATE TABLE MovieStar (  
    name CHAR(30) PRIMARY KEY,  
    address VARCHAR(255),  
    gender CHAR(1));
```

OR:

```
CREATE TABLE MovieStar (  
    name CHAR(30),  
    address VARCHAR(255),  
    gender CHAR(1)  
    PRIMARY KEY (name));
```

Primary Keys with Multiple Attributes

```
CREATE TABLE MovieStar (  
    name CHAR(30),  
    address VARCHAR(255),  
    gender CHAR(1),  
    PRIMARY KEY (name, address));
```

Other Keys

```
CREATE TABLE MovieStar (  
    name CHAR(30),  
    address VARCHAR(255),  
    phone CHAR(10) UNIQUE,  
    gender CHAR(1),  
    petName CHAR(50),  
    PRIMARY KEY (name),  
    UNIQUE (gender, petName));
```



Foreign Key Constraints

```
CREATE TABLE ActedIn (  
    Name CHAR(30) PRIMARY KEY,  
    MovieName CHAR(30)  
        REFERENCES Movies(MovieName),  
    Year INT);
```

Foreign key constraint doesn't prevent null, as it's necessary in self relation (If trigger to ensure insert, make FK null & warning).

Foreign Key Constraints

- OR

```
CREATE TABLE ActedIn (  
    Name CHAR(30) PRIMARY KEY,  
    MovieName CHAR(30),  
    Year INT,  
    FOREIGN KEY MovieName  
        REFERENCES Movies(MovieName)
```

- MovieName must be a PRIMARY KEY

How do we Maintain them?

- **Given a change to DB, there are several possible violations:**
 - Insert new tuple with bogus foreign key value
 - Update a tuple to a bogus foreign key value
 - Delete a tuple in the referenced table with the referenced foreign key value
 - Update a tuple in the referenced table that changes the referenced foreign key value

How to Maintain?

- **Recall, ActedIn has FK MovieName...**

Movies(MovieName, year)

(Fatal Attraction, 1987)

ActedIn(ActorName, MovieName)

(Michael Douglas, Fatal Attraction)

insert: (Rick Moranis, Strange Brew)

How to Maintain?

- **Policies for handling the change...**
 - Reject the update (default)
 - Cascade (example: cascading deletes)
 - Set NULL
- **Can set update and delete actions independently in CREATE TABLE**

MovieName CHAR(30)

REFERENCES Movies(MovieName))

ON DELETE SET NULL

ON UPDATE CASCADE